


# EdgeX Geneva 规则引擎介绍

## — 基于 EMQ X Kuiper 的超轻量级 IoT 流式数据处理


Apr, 2020

Rocky@emqx.io

杭州映云科技有限公司



# EMQ – 开源物联网消息接入与数据处理领导者



1

商业化开源软件

2

服务于 5G 时代的 IoT 产业

3

消息与流处理

4


全球 5000+ 企业用户

5

全球运营: 中国、美国硅谷、欧洲

# EMQ X Kuiper

- 将 Apache Flink、Spark 运行在边缘端！
- 原生二进制可运行包
- 基于 SQL 的业务逻辑实现方式
- 基于 Golang 的可扩展框架
  - 数据源
  - 数据目标 (Sinks)
  - 函数
- 规则管理的 API 与 CLI
- 基于 Apache 2.0 开源协议



项目地址: <https://github.com/emqx/kuiper>

# Kuiper 使用过程

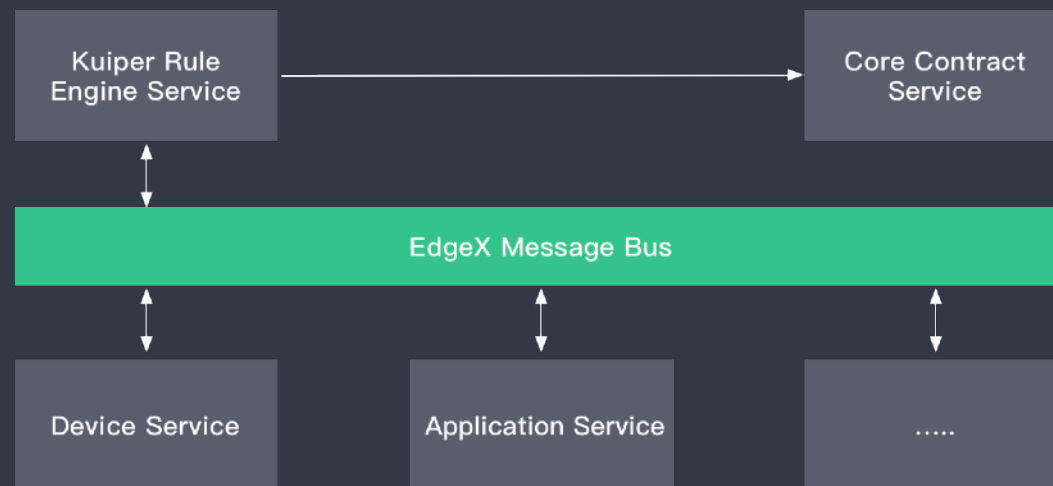
- 定义流
  - 类似于数据库中表的定义
- 定义并提交规则
  - 用 SQL 实现业务逻辑，并将运行结果发送到指定目标
  - 支持的 SQL
    - SELECT/FROM/WHERE/ORDER
    - JOIN/GROUP/HAVING
    - 4类时间窗口
    - 60+ SQL 函数
- 规则运行

```
# bin/cli create stream demo '(temperature float, humidity bigint)
WITH (FORMAT="JSON", DATASOURCE="devices/+/messages")'
```

```
{
  "sql": "SELECT avg(temperature) AS t_av, max(temperature) AS
t_max, min(temperature) AS t_min, COUNT(*) As t_count,
split_value(mqtt(topic), \"/\", 1) AS device_id FROM demo GROUP BY
device_id, TUMBLINGWINDOW(ss, 10)",
  "actions": [
    {
      "log": {}
    },
    {
      "mqtt": {
        "server": "ssl://xyz-ats.iot.us-east-1.amazonaws.com:8883",
        "topic": "devices/result",
        "qos": 1,
        "clientId": "demo_001",
        "certificationPath": "/var/aws/d3807d9fa5-certificate.pem",
        "privateKeyPath": "/var/aws/d3807d9fa5-private.pem.key"
      }
    }
  ]
}
```

# EdgeX 微服务之间的数据交互

- EdgeX 数据总线
  - 微服务之间的数据交流总线
  - 支持 ZeroMQ 与 MQTT
- 数据类型的处理
  - 从 Core Contract Service 中读取数据类型
  - 将数据转换为 Kuiper 的数据类型
  - EdgeX 中无需定义字段的类型，使用 schema-less 流定义
- 支持：Bool、String、Uint8、Uint16、Uint32、  
Uint64、Int8、Int16、Int32、Int64、Float32、  
Float64
- Binary, Array – 不支持



# EdgeX Event & Kuiper 数据映射

## 创建流

```
create stream events() WITH (FORMAT="JSON", TYPE="edgex")
```

## models.Event 数据样例

```
Events:
{Device: "demo", Created: 000, ...}

reading[0]: {Name: "Temperature", value: "30", "Created":123 ...}

reading[1]: {Name: "Humidity", value: "20", "Created":456 ...}
```

1) SELECT **temperature, humidity** FROM **events** WHERE **meta(device) = "demo"**

<b>temperature</b>	<b>humidity</b>
30	20

2) 使用 meta 函数的几个例子:

- meta(created): 000 //Get 'created' metadata from Event structure
- meta(temperature -> created): 123 //Get 'created' metadata from reading[0], key with 'temperature'
- meta(humidity -> created): 456 //Get 'created' metadata from reading[1], key with 'humidity'

# EdgeX 源使用例子

## 1) 使用 ZeroMQ 消息总线:

```
create stream events() WITH (FORMAT="JSON", TYPE="edgex")
```

etc/sources/edgex.yaml:

```
default:
  protocol: tcp
  server: edgex-app-service-configurable-rules
  port: 5566
  topic: events
  serviceServer: http://localhost:48080
```

## 2) 使用 MQTT 消息总线:

```
create stream demo () WITH (FORMAT="JSON", TYPE="edgex" Conf_key="mqtt_conf")
```

etc/sources/edgex.yaml:

```
mqtt_conf: #Conf_key
  protocol: tcp
  server: 127.0.0.1
  port: 1883
  topic: events
  type: mqtt
  optional:
    ClientId: "client1"
```

# EdgeX sink

- 往 EdgeX 消息总线发送符合规格的数据
- 元数据保留方式
  - 发布结果到 EdgeX 消息总线，而不保留原有的元数据：Kuiper 在此情况下作为 EdgeX 的一个单独微服务，它有自己的 device name。提供了属性 `deviceName`，该属性允许用户指定 Kuiper 的设备名称
  - 发布结果到 EdgeX 消息总线，并保留原有的元数据：在此情况下，Kuiper 更像是一个过滤器 – 将不关心的数据过滤掉，但是依然保留原有的数据。



# EdgeX sink 使用例子

```
{
  "id": "rule1",
  "sql": "SELECT temperature, humidity, meta(*) AS metadt FROM demo WHERE temperature = 72",
  "actions": [
    {
      "edgex": {
        "protocol": "tcp",
        "host": "*",
        "port": 5571,
        "topic": "application",
        "metadata": "metadt",
        "contentType": "application/json"
      }
    }
  ]
}
```

```
{
  "id": "rule1",
  "sql": "SELECT meta(*) AS edgex_meta, temperature, humidity, humidity*2 as h1 FROM demo WHERE temperature = 20",
  "actions": [
    {
      "edgex": {
        "protocol": "tcp",
        "host": "${mqtt_srv}",
        "port": 1883,
        "topic": "result",
        "type": "mqtt",
        "metadata": "edgex_meta",
        "contentType": "application/json",
        "optional": {
          "ClientId": "edgex_message_bus_001"
        }
      }
    },
    {
      "log": {}
    }
  ]
}
```

# Kuiper 性能 – 支持规则数目的测试

- 8000 规则 \* 0.1 消息/秒/规则, 共计的 TPS 为 800 条/秒
- 规则定义
  - 源: MQTT
  - SQL: select temperature from source where temperature > 20 (90% 数据被过滤)
  - 目标: 日志
- 配置
  - AWS: 2core \* 4GB
  - Ubuntu
- 资源使用
  - Memory: 89% ~ 72%; 0.4 MB/rule
  - CPU: 25%

# 性能基准测试

- 配置
  - AWS t2.micro( 1 Core \* 1 GB)
  - Ubuntu 18.04
- 场景
  - 一个 Go 应用用于往 ZeroMQ 消息总线发送数据，格式如下：

```
{ "device": "demo",  
  "readings": [  
    { "device": "Temperature device", "name": "Temperature", "value": "10"},  
    { "device": "Humidity device", "name": "Humidity", "value": "15"}  
  ]  
}
```
  - 规则如下
    - A nop sink – 忽略发送到该 sink 的所有数据
    - 过滤 90% 数据
- 结果
  - 11434 message/秒
  - 75% CPU (user + sys)
  - 内存: 32MB

```
top - 04:26:10 up 1 day, 22:11, 3 users, load average: 1.03, 0.44, 0.16
Tasks: 98 total, 1 running, 61 sleeping, 0 stopped, 0 zombie
%Cpu(s): 56.6 us, 11.2 sy, 0.0 ni, 25.3 id, 0.0 wa, 0.0 hi, 6.4 si, 0.4 st
KiB Mem : 1007300 total, 369696 free, 175060 used, 462544 buff/cache
KiB Swap: 0 total, 0 free, 0 used. 681012 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
8259	ubuntu	20	0	780964	32280	4804	S	74.8	3.2	3:32.95	server
8583	ubuntu	20	0	44580	4212	3524	R	0.3	0.4	0:00.68	top

```
{  
  "sql": "SELECT * from demo where temperature>50",  
  "actions": [  
    {  
      "nop": {  
        "log": false  
      }  
    }  
  ]  
}
```

# 测试状态截屏

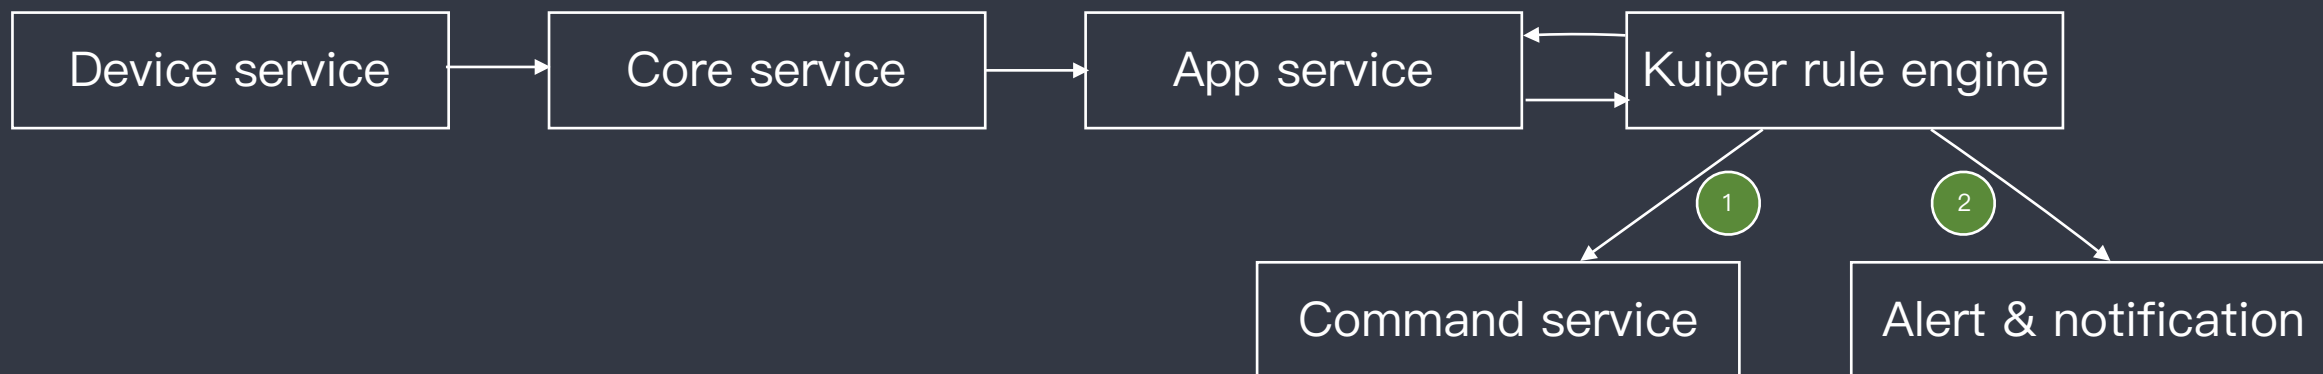
```
ubuntu@ip-172-31-1-144:~$ ./pub 1000000
elapsed 174.924363s
```

```
ubuntu@ip-172-31-5-85:/tmp/kuiper-master/_build/kuiper--linux-x86_64$ bin/cli getstatus rule rule1
Connecting to 127.0.0.1:20498...
{
  "source_demo_0_records_in_total": 1000000,
  "source_demo_0_records_out_total": 1000000,
  "source_demo_0_exceptions_total": 0,
  "source_demo_0_process_latency_ms": 0,
  "source_demo_0_buffer_length": 0,
  "source_demo_0_last_invocation": "2020-04-10T04:26:15.51329",
  "op_preprocessor_demo_0_records_in_total": 1000000,
  "op_preprocessor_demo_0_records_out_total": 1000000,
  "op_preprocessor_demo_0_exceptions_total": 0,
  "op_preprocessor_demo_0_process_latency_ms": 0,
  "op_preprocessor_demo_0_buffer_length": 0,
  "op_preprocessor_demo_0_last_invocation": "2020-04-10T04:26:15.513371",
  "op_filter_0_records_in_total": 1000000,
  "op_filter_0_records_out_total": 100000,
  "op_filter_0_exceptions_total": 0,
  "op_filter_0_process_latency_ms": 0,
  "op_filter_0_buffer_length": 0,
  "op_filter_0_last_invocation": "2020-04-10T04:26:15.513391",
  "op_project_0_records_in_total": 100000,
  "op_project_0_records_out_total": 100000,
  "op_project_0_exceptions_total": 0,
  "op_project_0_process_latency_ms": 0,
  "op_project_0_buffer_length": 0,
  "op_project_0_last_invocation": "2020-04-10T04:26:15.513468",
  "sink_sink_nop_0_records_in_total": 100000,
  "sink_sink_nop_0_records_out_total": 100000,
  "sink_sink_nop_0_exceptions_total": 0,
  "sink_sink_nop_0_process_latency_ms": 0,
  "sink_sink_nop_0_buffer_length": 1,
  "sink_sink_nop_0_last_invocation": "2020-04-10T04:26:15.513501"
}
```

```
ubuntu@ip-172-31-5-85:/tmp/kuiper-master/_build/kuiper--linux-x86_64$ bin/cli getstatus rule rule2
Connecting to 127.0.0.1:20498...
{
  "source_demo_0_records_in_total": 1000000,
  "source_demo_0_records_out_total": 1000000,
  "source_demo_0_exceptions_total": 0,
  "source_demo_0_process_latency_ms": 0,
  "source_demo_0_buffer_length": 0,
  "source_demo_0_last_invocation": "2020-04-10T04:26:15.514621",
  "op_preprocessor_demo_0_records_in_total": 1000000,
  "op_preprocessor_demo_0_records_out_total": 1000000,
  "op_preprocessor_demo_0_exceptions_total": 0,
  "op_preprocessor_demo_0_process_latency_ms": 0,
  "op_preprocessor_demo_0_buffer_length": 0,
  "op_preprocessor_demo_0_last_invocation": "2020-04-10T04:26:15.514631",
  "op_filter_0_records_in_total": 1000000,
  "op_filter_0_records_out_total": 100000,
  "op_filter_0_exceptions_total": 0,
  "op_filter_0_process_latency_ms": 0,
  "op_filter_0_buffer_length": 0,
  "op_filter_0_last_invocation": "2020-04-10T04:26:15.514635",
  "op_project_0_records_in_total": 100000,
  "op_project_0_records_out_total": 100000,
  "op_project_0_exceptions_total": 0,
  "op_project_0_process_latency_ms": 0,
  "op_project_0_buffer_length": 0,
  "op_project_0_last_invocation": "2020-04-10T04:26:15.514639",
  "sink_sink_nop_0_records_in_total": 100000,
  "sink_sink_nop_0_records_out_total": 100000,
  "sink_sink_nop_0_exceptions_total": 0,
  "sink_sink_nop_0_process_latency_ms": 0,
  "sink_sink_nop_0_buffer_length": 1,
  "sink_sink_nop_0_last_invocation": "2020-04-10T04:26:15.514652"
}
```

# EdgeX Kuiper 规则使用场景

- Kuiper 规则引擎的可能输出
  - Command service: send control command (rest API)
  - Alert & notification: send alert & notification (rest API)



# 2020 规划

- EdgeX Hanoi Release
  - 数组类型支持
  - Sink 数据发送的时候提供更多的灵活性 (Golang data template)
  - UI 提升
- 与 KubeEdge 集成
- 我们非常乐意看到来自于社区的新需求!
- 欢迎加入我们的项目!
  
- 规划链接: <https://github.com/emqx/kuiper/projects/1>

# 文档

---

- 项目网址: <https://github.com/emqx/kuiper>
- Docker: <https://hub.docker.com/r/emqx/kuiper>
- EdgeX Kuiper 使用教程: [https://github.com/emqx/kuiper/blob/master/docs/zh\\_CN/edgex/edgex\\_rule\\_engine\\_tutorial.md](https://github.com/emqx/kuiper/blob/master/docs/zh_CN/edgex/edgex_rule_engine_tutorial.md)
- 插件开发教程: [https://github.com/emqx/kuiper/blob/master/docs/zh\\_CN/plugins/plugins\\_tutorial.md](https://github.com/emqx/kuiper/blob/master/docs/zh_CN/plugins/plugins_tutorial.md)

A series of thin, green, wavy lines that flow horizontally across the middle of the slide, creating a sense of motion and depth.

# Thank You

[contact@emqx.io](mailto:contact@emqx.io)

